Bayes Decision Theory

Abraham Korh

Contents

Introduction

In this assignment we'll construct a Naive Bayes Classifier, using several ideas, but focus on to very important ideas from bayesian probability theory. The first being **independence**, which in layman terms, essentially refers to the property of a probability for one thing occurring is unrelated and unaffected of the probability of another thing occurring. The second being **Bayes Theorem**, which essentially describes the relationship between single probabilities and conditional probabilities.

Prerequisites

Independence

Exclusivity

Independence depends on the fact that the outcomes or possibilities of whatever space, universe or phenomenon we're interested in, are mutually exclusive. Essentially, in such a universe, any sample maps to a single outcome among many possible outcomes. Formally, we can describe exclusiveness using set theory. A series of events $(A_1, A_2, ..., A_n)$ are mutually exclusive if their intersection is equal to an empty set.

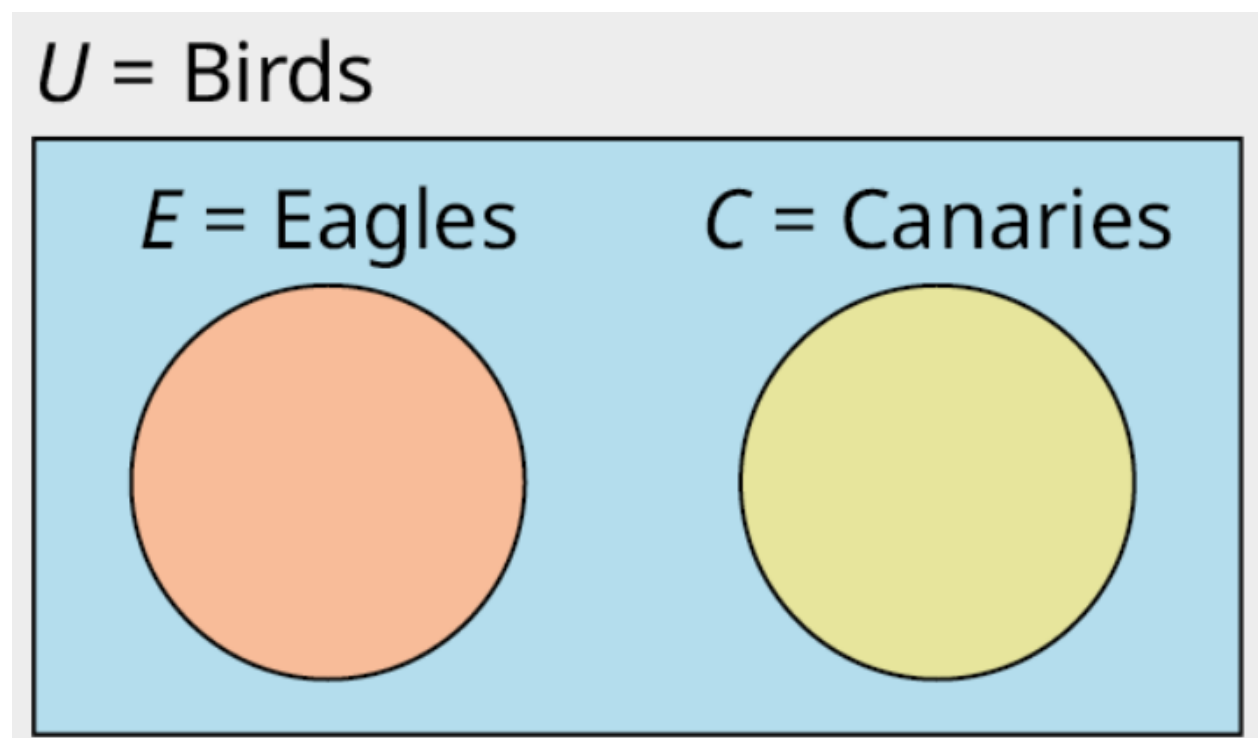$$A_i \cap A_j = \emptyset \quad \text{for all } i \neq j$$

Fig. 1: Illustration of mutual exclusivity.

Back to Independence

The concept of independence only holds when the various outcomes are mutually exclusive. In that case, independence allows us to express the probability of a conjunction $P(A_1 \cap A_2 \cap ... \cap A_n)$ as the product

of their single probabilities

$$P(\ A_1\ \cap\ A_2\ \cap\ ...\ \cap\ A_n\ ) = \prod_{i \in I} P(A_i)$$

Bayes Theorem

Conditional Probabilities

The strength of Bayes Theorem lies in its ability to update the probability estimate for an event as more information becomes available. This is achieved through the use of conditional probabilities. Conditional probability is the probability of an event occurring given that another event has already occurred.

The formula for conditional probability is given by:

$$P(A \mid B) = \frac{P(A\ \cap\ B)}{P(B)}$$

where $P(A \mid B)$ is the probability of event A occurring given that B has occurred, and $P(A\ \cap\ B)$ is the probability of both events A and B occurring.

Back to Bayes Theorem

Bayes Theorem is a way of finding a probability when we know certain other probabilities. The classic equation for Bayes Theorem is:

$$P(A \mid B) = \frac{P(B \mid A) \cdot P(A)}{P(B)}$$

This theorem allows us to update our prior beliefs with new evidence.

Total Probability

The Law of Total Probability is a fundamental rule relating marginal probabilities to conditional probabilities. It states that the probability of an event can be found by considering all possible ways that the event can occur. For a set of mutually exclusive and exhaustive events B_i, the law is given by:

$$P(A) = \sum_{i \in \mathbb{I}} P(B_i \mid A) \cdot P(B_i)$$

This theorem is particularly useful in scenarios where the probability of an event A is not directly known, but can be broken down into several mutually exclusive events B_i.

Constructing Our Naive Bayes Classifier

Our Starting Point

Given a set of features $\mathbf{X} = (x_1, x_2, ..., x_i)$ and set of target classes $\mathbf{Y} = (y_1, y_2, ..., y_k)$, we can begin constructing a Naive Bayes classifier. We begin by defining our decision rule by the following expression

$$\hat{y} = \arg\max_y P(y|\mathbf{X})$$

Which can then be expanded to

$$\hat{y} = \arg\max_y P(y|\mathbf{X}) \tag{1}$$

$$= \arg\max_y \frac{P(\mathbf{X}|y) \cdot P(y)}{P(\mathbf{X})} \tag{2}$$

Putting it together

Feature Independence

Assuming each feature $x_i \in \mathbf{X}$ is independent of the remaining features, then the probability of the conjunction $P(\mathbf{X})$ can be expressed as

$$P(\mathbf{X}) = P(x_1, x_2, ..., x_i) = \prod_{i \in I} P(x_i)$$

And as a result the conditional probability of the conjunction $P(\mathbf{X}|y)$ can be expressed as

$$P(\mathbf{X}|y) = P(x_1, x_2, ..., x_i) = \prod_{i \in I} P(x_i|y)$$

Applying total probability

Now we turn our attention to ther term in the numerator $P(\mathbf{X})$ which can be expressed as a sum of over the product of conditional and single probabilities

$$P(\mathbf{X}) = \sum_{k \in \mathbb{K}} P(y_k \mid \mathbf{X}) \cdot P(y_i)$$

The Decision Rule

We can now substitute various components in order to arrive at the full expression

$$\hat{y} = \arg\max_y \left(\sum_{k \in \mathbb{K}} P(y_k \mid \mathbf{X}) \cdot P(y_i) \right)^{-1} \cdot P(y) \cdot \prod_{i \in I} P(x_i|y) \tag{3}$$

Finishing Notes

In practice, we compute values for $P(y)$, $P(y|\mathbf{X})$, and $P(x_i|y)$ directly from the dataset, and refer to them as priors. This then gives us everything we need in order to make use of our decision rule.

References

Figure 1: Understanding Venn Diagrams